Adverse Selection in the AI Data Commons

Evidence from News and Media Sites

Kai Zhu

Bocconi University

Digital Economy Workshop 2026

2026-03-13

Adverse Selection in AI Data Commons

Hello everyone. Thank you for having me. Today I'm presenting work on adverse selection in the AI data commons, focusing on news and media sites.

The central question: when websites can opt out of AI crawling, who opts out—and what does that mean for the quality of what remains?

The AI Training Data Problem

- Generative AI derives its power from training on **high-quality web content** (Gunasekar et al. 2023; Longpre et al. 2023)

The AI Training Data Problem
Generative AI derives its power from training on high-quality web content (Gunasekar et al. 2023; Longpre et al. 2023)

Here is the core problem. Training data quality drives model quality—Gunasekar and colleagues showed that curated, high-quality data can match models trained on far larger but noisier corpora. The most valuable content comes from producers who invest in editorial standards and original reporting.

Yet AI companies use this content without compensation. There are a handful of deals—OpenAI paid Reddit 60 million dollars, signed a licensing agreement with the AP—but these are rare exceptions.

Without a market, publishers face a binary choice: allow the crawler or block it via robots.txt. So the question becomes: who opts out? If the highest-quality producers exit first, we get a classic lemons problem. The training corpus degrades not because many sites leave, but because the *best* sites leave. That is exactly what we find.

The AI Training Data Problem

- Generative AI derives its power from training on **high-quality web content**
  (Gunasekar et al. 2023; Longpre et al. 2023)
- Yet AI companies use this content **without compensation**
  Rare exceptions:
  - OpenAI–Reddit $60M deal
  - OpenAI–Associated Press licensing agreement

The AI Training Data Problem
Generative AI derives its power from training on high-quality web content
(Gunasekar et al. 2023; Longpre et al. 2023)
Yet AI companies use this content without compensation
Rare exceptions:
OpenAI–Reddit $60M deal
OpenAI–Associated Press licensing agreement

Here is the core problem. Training data quality drives model quality—Gunasekar and colleagues showed that curated, high-quality data can match models trained on far larger but noisier corpora. The most valuable content comes from producers who invest in editorial standards and original reporting.

Yet AI companies use this content without compensation. There are a handful of deals—OpenAI paid Reddit 60 million dollars, signed a licensing agreement with the AP—but these are rare exceptions.

Without a market, publishers face a binary choice: allow the crawler or block it via robots.txt. So the question becomes: who opts out? If the highest-quality producers exit first, we get a classic lemons problem. The training corpus degrades not because many sites leave, but because the *best* sites leave. That is exactly what we find.

The AI Training Data Problem

- Generative AI derives its power from training on **high-quality web content**
  (Gunasekar et al. 2023; Longpre et al. 2023)
- Yet AI companies use this content **without compensation**
  Rare exceptions:
  - OpenAI–Reddit $60M deal
  - OpenAI–Associated Press licensing agreement
- No established market: producers' only option is **binary opt-out** via `robots.txt`
  (Kim et al. 2025; Chang & He 2025)

The AI Training Data Problem
Generative AI derives its power from training on high-quality web content
(Gunasekar et al. 2023; Longpre et al. 2023)
Yet AI companies use this content without compensation
Rare exceptions:
OpenAI–Reddit $60M deal
OpenAI–Associated Press licensing agreement
No established market: producers' only option is binary opt-out via robots.txt
(Kim et al. 2025; Chang & He 2025)

Here is the core problem. Training data quality drives model quality—Gunasekar and colleagues showed that curated, high-quality data can match models trained on far larger but noisier corpora. The most valuable content comes from producers who invest in editorial standards and original reporting.

Yet AI companies use this content without compensation. There are a handful of deals—OpenAI paid Reddit 60 million dollars, signed a licensing agreement with the AP—but these are rare exceptions.

Without a market, publishers face a binary choice: allow the crawler or block it via robots.txt. So the question becomes: who opts out? If the highest-quality producers exit first, we get a classic lemons problem. The training corpus degrades not because many sites leave, but because the *best* sites leave. That is exactly what we find.

The AI Training Data Problem

- Generative AI derives its power from training on **high-quality web content**
  (Gunasekar et al. 2023; Longpre et al. 2023)
- Yet AI companies use this content **without compensation**
  Rare exceptions:
  - OpenAI–Reddit $60M deal
  - OpenAI–Associated Press licensing agreement
- No established market: producers' only option is **binary opt-out** via `robots.txt`
  (Kim et al. 2025; Chang & He 2025)

Who opts out? If the highest-quality producers exit first, the result is a **lemons problem** (Akerlof, 1970)—and the remaining AI data commons systematically degrades.

The AI Training Data Problem
Generative AI derives its power from training on high-quality web content (Gunasekar et al. 2023; Longpre et al. 2023)
Yet AI companies use this content without compensation
Rare exceptions:
OpenAI–Reddit $60M deal
OpenAI–Associated Press licensing agreement
No established market: producers' only option is binary opt-out via robots.txt (Kim et al. 2025; Chang & He 2025)
Who opts out? If the highest-quality producers exit first, the result is a lemons problem (Akerlof, 1970)—and the remaining AI data commons systematically degrades.

Here is the core problem. Training data quality drives model quality—Gunasekar and colleagues showed that curated, high-quality data can match models trained on far larger but noisier corpora. The most valuable content comes from producers who invest in editorial standards and original reporting.

Yet AI companies use this content without compensation. There are a handful of deals—OpenAI paid Reddit 60 million dollars, signed a licensing agreement with the AP—but these are rare exceptions.

Without a market, publishers face a binary choice: allow the crawler or block it via robots.txt. So the question becomes: who opts out? If the highest-quality producers exit first, we get a classic lemons problem. The training corpus degrades not because many sites leave, but because the *best* sites leave. That is exactly what we find.

Research Questions & Preview of Findings

RQ: Does the market for AI training data exhibit adverse selection?

Key findings:

1. **Quality–blocking gradient**
   High-factual media block at **6**× the rate of low-factual sources
2. **Strategic targeting**
   High-quality sites selectively block training bots over search bots
3. **Misinformation asymmetry**
   Conspiracy, propaganda, pseudoscience stay open for AI crawlers
4. **Ideological sorting**
   Centrist outlets block most; both left and right extremes block least

⇒ Evidence from **9,600 news and media sites** with two complementary quality measures: expert ratings and web-graph metrics

Research Questions & Preview of Findings
RQ: Does the market for AI training data exhibit adverse selection?
Key findings:
1. Quality–blocking gradient
High-factual media block at 6× the rate of low-factual sources
2. Strategic targeting
High-quality sites selectively block training bots over search bots
3. Misinformation asymmetry
Conspiracy, propaganda, pseudoscience stay open for AI crawlers
4. Ideological sorting
Centrist outlets block most; both left and right extremes block least
⇒ Evidence from 9,600 news and media sites with two complementary quality measures: expert ratings and web-graph metrics

Our central research question is whether the AI training data market exhibits adverse selection. And the answer is a clear yes.

Let me preview the key findings. We see a steep quality–blocking gradient—high-factual media block at six times the rate of low-factual sources. This blocking is strategic: sites target training bots while keeping search bots that drive traffic. The misinformation dimension is particularly striking—conspiracy, propaganda, and pseudoscience sources remain freely available while credible outlets opt out. And there is an ideological pattern too: centrist outlets block most, while both political extremes stay in the training pool.

We test all of this using 9,600 news sites with two complementary quality measures: expert ratings from Media Bias/Fact Check and algorithmic web-graph metrics. Let me walk you through the setup.

Institutional Background

robots.txt

- Voluntary protocol since 1994
- Sites specify which bots to block
- Advisory, not legally enforced
- Can target *specific* AI crawlers

AI Crawlers

- **Training**: GPTBot, Google-Extended, ClaudeBot
  Extract content to build models
- **Search**: ChatGPT-User, PerplexityBot, Applebot-Extended
  Retrieve content to answer queries — can drive traffic back

Key events:

Aug 2023	Dec 2023	Jul 2024
GPTBot launched	NYT lawsuit	Cloudflare AI block

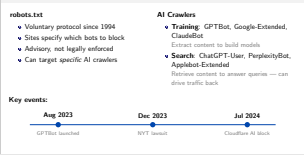
Institutional Background
robots.txt
Voluntary protocol since 1994
Sites specify which bots to block
Advisory, not legally enforced
Can target specific AI crawlers
AI Crawlers
Training: GPTBot, Google-Extended, ClaudeBot
Extract content to build models
Search: ChatGPT-User, PerplexityBot, Applebot-Extended
Retrieve content to answer queries — can drive traffic back
Key events:
Aug 2023
GPTBot launched
Dec 2023
NYT lawsuit
Jul 2024
Cloudflare AI block

Let me give you the institutional context. Robots.txt is a 30-year-old voluntary protocol that lets websites tell crawlers what not to access. It is advisory—not legally enforced—but sites can target specific AI crawlers by name.

The distinction between training and search bots is central to our analysis. Training bots like GPTBot extract content to build models—they take value from the site. Search bots like ChatGPT-User retrieve content to answer queries and can drive traffic back. Producers face very different incentives for each type.

Three events structure our timeline. GPTBot's launch in August 2023 made AI-specific opt-out salient for the first time. The NYT lawsuit in December 2023 raised the legal stakes. And Cloudflare's one-click blocking tool in July 2024 dramatically lowered the cost of opting out. With that context, let me explain why we focus specifically on media sites.

Why Media? Data & Sample

Why media & news?

- News content shapes public discourse and how AI models handle factual claims
- Expert quality ratings exist—enabling direct measurement of adverse selection
- Center of active policy debate (NYT v. OpenAI, EU AI Act)
- **High stakes**: if credible outlets exit but misinformation stays, AI inherits a distorted information environment

Sample

- **9,611** media/news sites (cross-section, February 2026)
- **7,002** sites $\times$ 45 months panel (August 2022–February 2026)
- Source: Media Bias/Fact Check (MBFC)
- robots.txt scraped monthly via HTTP Archive

Why Media? Data & Sample
Why media & news?
News content shapes public discourse and how AI models handle factual claims
Expert quality ratings exist—enabling direct measurement of adverse selection
Center of active policy debate (NYT v. OpenAI, EU AI Act)
High stakes: if credible outlets exit but misinformation stays, AI inherits a distorted information environment
Sample
9,611 media/news sites (cross-section, February 2026)
7,002 sites × 45 months panel (August 2022–February 2026)
Source: Media Bias/Fact Check (MBFC)
robots.txt scraped monthly via HTTP Archive

Why focus specifically on media and news? This sector is uniquely suited for studying adverse selection in AI training data. News content shapes public discourse and directly influences how AI models handle factual claims. Crucially, expert quality ratings exist for thousands of media sites, which lets us measure adverse selection directly rather than relying on proxies. And the stakes are high: if credible outlets exit while misinformation stays, AI models inherit a fundamentally distorted information environment.

Our sample comes from Media Bias/Fact Check, which provides expert ratings on factual reporting, credibility, and political orientation. We have 9,600 sites in cross-section and a panel of 7,000 sites tracked monthly over 45 months. Robots.txt files come from HTTP Archive. Now let me show how we measure quality.

Quality Measures: Two Complementary Indices

Content quality index *("inside-out")*

Expert-rated: what the site *produces*

- MBFC factual reporting (1–6)
- MBFC credibility rating (1–4)
- Wikipedia notable (binary)

→ PCA first component, standardized

Domain quality index *("outside-in")*

Algorithmic: how the *web* treats the site

- Trust Flow (Majestic, 0–100)
- Open PageRank (web centrality)
- Tranco rank (aggregated popularity)

→ PCA first component, standardized

Quality Measures: Two Complementary Indices
Content quality index ("inside-out")
Expert-rated: what the site produces
MBFC factual reporting (1–6)
MBFC credibility rating (1–4)
Wikipedia notable (binary)
→ PCA first component, standardized
Domain quality index ("outside-in")
Algorithmic: how the web treats the site
Trust Flow (Majestic, 0–100)
Open PageRank (web centrality)
Tranco rank (aggregated popularity)
→ PCA first component, standardized

We construct two composite quality indices via PCA, each capturing a distinct dimension of quality.

The content quality index is an inside-out measure based on expert ratings: MBFC factual reporting on a six-point scale, credibility on a four-point scale, and whether the site has a Wikipedia article. These reflect what the site produces, as assessed by trained human analysts.

The domain quality index is an outside-in measure based on algorithmic signals: Trust Flow captures link-graph authority, Open PageRank captures web centrality, and Tranco captures aggregated traffic. These reflect how the broader web treats the site.

Both are standardized to mean zero and unit variance. The key result is that both indices tell the same story—let me show you.

Result 1: Best Content Most Likely to Leave the Training Pool

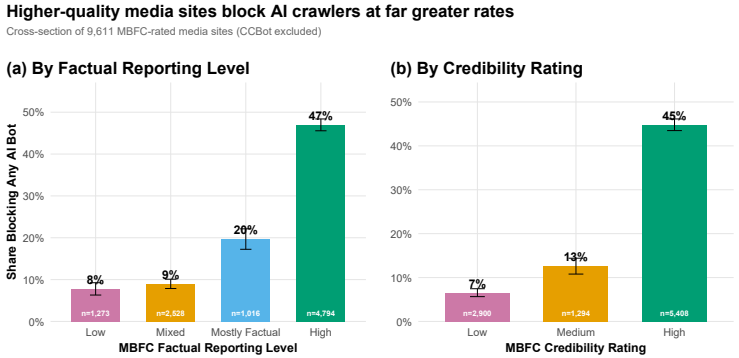

Textbook adverse selection:

- Low-factual sites: ∼**8%** block; high-factual sites: ∼**47%** block—a **6**× gap
- Monotonically increasing across both factual reporting and credibility
  Consistent across all four quality measures: factual, credibility, content quality index, domain quality index

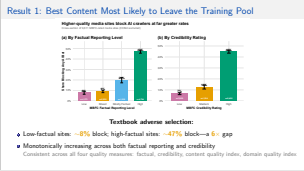

Here is our headline finding. The left panel plots blocking rates by MBFC factual reporting; the right panel by credibility rating. In both cases, blocking rises monotonically with quality.

The numbers are stark. Only about 8 percent of low-factual sources block AI training bots, compared to 47 percent of high-factual sources—a six-fold gap. Pause on that for a moment: nearly half of the most factual media sources have already opted out of AI training.

This gradient is not driven by a single measure. It is consistent across all four quality indicators we construct. The pattern is textbook adverse selection: the best content is the most likely to leave the training pool. And the consequences go beyond average quality declining.

Result 2: Misinformation Stays Freely Available for AI Training

(a) Credible vs. Questionable

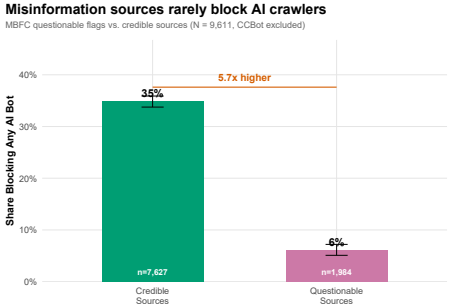

(b) By Misinformation Flag

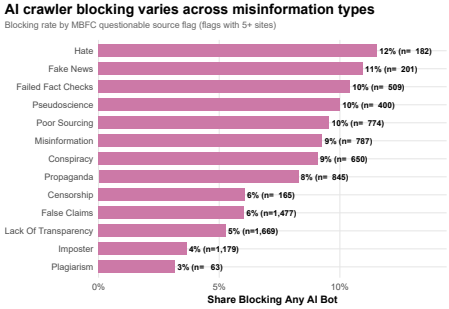

Opt-out skews the corpus toward misinformation:

- Credible outlets block at **5.7**× the rate of questionable sources (35% vs. 6%)
- Conspiracy, pseudoscience, propaganda sources: all below 10%
- As credible sources exit, the remaining corpus tilts toward misinformation

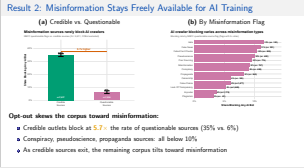

The consequences become clearest when we look at misinformation specifically. Panel (a) shows the headline comparison: credible outlets block training bots at nearly six times the rate of MBFC-flagged questionable sources—35 percent versus just 6 percent.

Panel (b) breaks this down by specific misinformation flag. Across the board—conspiracy, pseudoscience, propaganda, fake news—blocking rates are in the single digits. None exceed 10 percent.

So this is not just about average quality declining. The composition of what remains is actively shifting toward unreliable content. High-quality journalism opts out; misinformation stays freely available. The training corpus is not just getting worse—it is getting systematically skewed. There is also an ideological dimension to this sorting.

Result 3: AI Training Data Is Becoming Ideologically Sorted

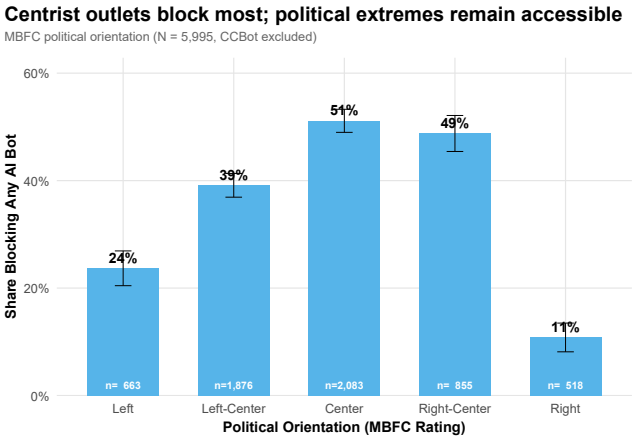

Political extremes stay in the training pool:

- Center: **51%** block
- Left: 24% block
- Right: only **11%** block

⇒ Inverted-U: centrist voices exit, extremes remain

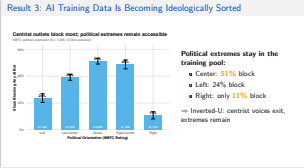

Beyond quality and misinformation, we find a striking ideological pattern. Center and center-right outlets block at the highest rates—around 49 to 51 percent. Left-leaning outlets block at 24 percent, and right-leaning outlets at just 11 percent.

The pattern forms an inverted-U: centrist voices disproportionately opt out, while both political extremes remain in the training pool. This means AI training data is becoming ideologically sorted—not just lower quality, but systematically skewed away from the political center.

A natural question arises: is this just because quality differs across the political spectrum? Let me show the regression evidence.

Result 3: Ideological Gaps Persist After Quality Controls

	(1) Bias	(2) +Dom. Q	(3) +Cont. Q	(4) Full
Left	−.354***	−.352***	−.234***	−.230***
	(.024)	(.024)	(.026)	(.026)
Left-Ctr	−.173***	−.177***	−.159***	−.163***
	(.017)	(.017)	(.017)	(.017)
Right-Ctr	−.101***	−.102***	−.045**	−.045**
	(.022)	(.022)	(.022)	(.022)
Right	−.483***	−.467***	−.206***	−.186***
	(.028)	(.028)	(.035)	(.035)
Cont. Q			.234***	.236***
			(.018)	(.018)
N	4,650	4,650	4,650	4,650
R^2	.085	.091	.117	.125
Cat. FE				X

Ref: Center. * $p < .1$, ** $p < .05$, *** $p < .01$

Quality explains part, but not all:

- Right: **−19 pp** vs. Center after full controls
- Left: **−23 pp** vs. Center after full controls
- Content quality is a strong predictor (+24 pp per SD)

⇒ Political orientation has an **independent** effect beyond quality

Result 3: Ideological Gaps Persist After Quality Controls

	(1) Bias	(2) +Dom. Q	(3) +Cont. Q	(4) Full
Left	−.354***	−.352***	−.234***	−.230***
	(.024)	(.024)	(.026)	(.026)
Left-Ctr	−.173***	−.177***	−.159***	−.163***
	(.017)	(.017)	(.017)	(.017)
Right-Ctr	−.101***	−.102***	−.045**	−.045**
	(.022)	(.022)	(.022)	(.022)
Right	−.483***	−.467***	−.206***	−.186***
	(.028)	(.028)	(.035)	(.035)
Cont. Q			.234***	.236***
			(.018)	(.018)
N	4,650	4,650	4,650	4,650
R^2	.085	.091	.117	.125
Cat. FE				X

Ref: Center. * $p < .1$, ** $p < .05$, *** $p < .01$

Quality explains part, but not all:
- Right: −19 pp vs. Center after full controls
- Left: −23 pp vs. Center after full controls
- Content quality is a strong predictor (+24 pp per SD)

⇒ Political orientation has an independent effect beyond quality

So is the ideological pattern just driven by quality differences? This table explicitly controls for quality.

Column 1 shows raw bias effects relative to center: right-leaning outlets are 48 percentage points less likely to block. Column 4 adds domain quality, content quality, and category fixed effects. The right-wing coefficient shrinks to minus 19 percentage points—quality explains about 60 percent of the gap—but a large, highly significant residual remains. Similarly, left-leaning outlets remain 23 percentage points below center after full controls.

The takeaway: political orientation has an independent effect on blocking behavior beyond what quality predicts. AI training data is becoming ideologically sorted by design, not merely as a byproduct of quality selection. Now let me turn to the temporal dynamics.

Result 4: Quality–Blocking Gradient Amplified After GPTBot

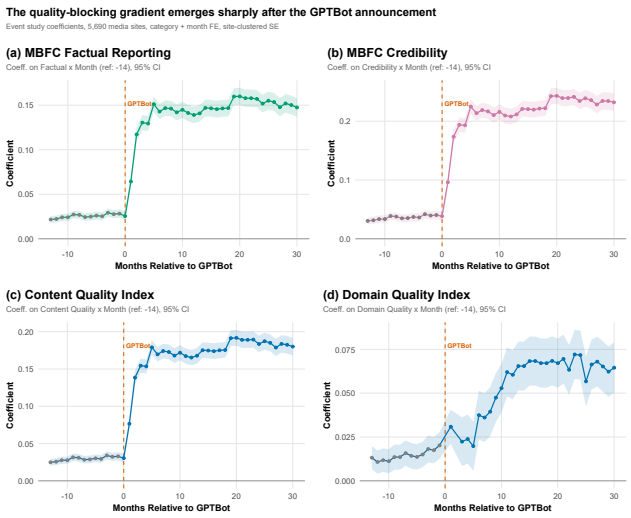

When did adverse selection emerge?

- Pre-period: small, flat quality–blocking gradient
- Aug 2023 (GPTBot): **sharp break**—coefficients jump **5–6**×
- Gradient **keeps widening**: social learning and cascading adoption

⇒ Not a pre-existing pattern—GPTBot made opt-out salient and actionable

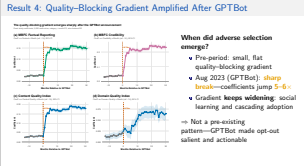

When exactly did this adverse selection emerge? This event study pins down the timing. Before August 2023, the quality–blocking gradient is small and flat—there is no meaningful adverse selection in the pre-period.

GPTBot's launch changed that. It was the first major AI crawler to announce itself publicly, making AI-specific opt-out salient and actionable for the first time. The break is sharp—post-period coefficients are five to six times larger. And critically, the gradient keeps widening over time, suggesting social learning and cascading adoption among publishers.

This pattern is consistent across all four quality measures. It raises a natural question: is the blocking strategic, or are sites simply blocking everything?

Result 5: Adverse Selection Concentrated in Training Data

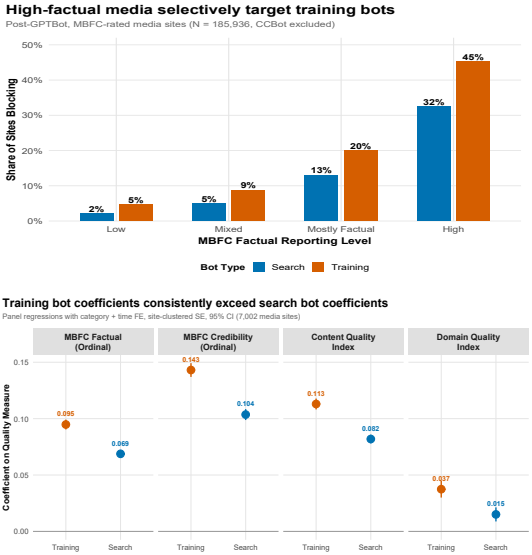

Blocking is strategic, not indiscriminate:

- Quality gap larger for training bots than search bots across factual levels
- Regression coefficients 2–3× larger for training bots across all quality measures
- Robust across cross-section, panel, site + time FE

⇒ Producers block value *extraction*; permit value *creation*

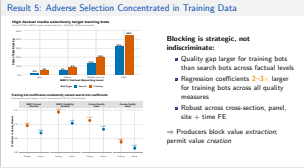

The answer is clear: blocking is highly strategic. High-quality sites disproportionately block training bots—the ones that extract content to build models—while permitting search bots that can drive traffic back.

The top panel shows the descriptive gap: across every factual reporting level, the quality gradient is steeper for training bots than for search bots. The bottom panel confirms this in regressions. The quality coefficient for blocking is two to three times larger for training bots than for search bots, and this holds across every quality measure—factual reporting, credibility, content quality index.

These results are robust across cross-section and panel specifications with site and time fixed effects. Producers clearly understand the distinction and act on it. So what does all of this mean in aggregate?

Result 6: Counterfactual Simulation — Selection, Not Volume, Drives Quality Loss

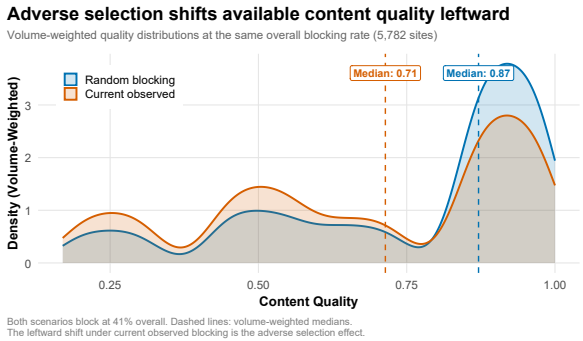

Decomposing the quality decline:

- **Counterfactual**: same blocking rate: quality-correlated blocking vs random selection
- Adverse selection accounts for the **entire** quality decline, not blocking per se
- Projections through **2028**: degradation compounds to **18–32%** quality decrease

⇒ Quality loss driven by *selection*, not by volume of opt-outs

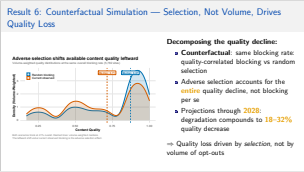

To quantify the aggregate impact, we run counterfactual simulations. We compare the actual training corpus—where blocking is quality-correlated—against a counterfactual where the same number of sites block but selection is random. The gap between these two scenarios isolates the pure effect of adverse selection.

The result is striking: adverse selection accounts for essentially the entire quality decline. If the same number of sites blocked randomly, the corpus would barely degrade. The problem is not that sites are opting out—it is which sites are opting out.

Projecting forward through 2028 using observed trends, the degradation compounds substantially—an 18 to 32 percent quality decrease depending on the scenario. This has direct policy implications.

Policy Implications

1. **Licensing frameworks**: Replace binary opt-out with a market for data licenses
   Compensate quality producers to stay in the commons
2. **Collective action**: Lower transaction costs via industry-wide standards
   cf. ASCAP for music — centralized licensing reduces bilateral negotiation costs
3. **Regulatory design**: Opt-out regimes must account for adverse selection
   EU AI Act, copyright reform

   Opt-out regimes by design produce a lemons equilibrium

Policy Implications
1. Licensing frameworks: Replace binary opt-out with a market for data licenses
Compensate quality producers to stay in the commons
2. Collective action: Lower transaction costs via industry-wide standards
cf. ASCAP for music — centralized licensing reduces bilateral negotiation costs
3. Regulatory design: Opt-out regimes must account for adverse selection
EU AI Act, copyright reform
Opt-out regimes by design produce a lemons equilibrium

So what can be done? Three policy implications follow from our findings.
First, we need to move from binary opt-out to structured licensing frameworks. Give quality producers compensation for staying in the commons, rather than forcing a take-it-or-leave-it choice.
Second, collective action mechanisms could reduce transaction costs. Think ASCAP for web content—industry-wide standards that let publishers license collectively, rather than negotiating bilateral deals one at a time.
Third, and most importantly for regulators: opt-out regimes inherently produce adverse selection. The EU AI Act and ongoing copyright reforms need to account for this. If the mechanism is opt-out, the equilibrium is lemons—by design, not by accident.

Conclusion

Three contributions:

1. **Economics of data markets**: First evidence of adverse selection in AI training data
   No price mechanism → best content exits first → lemons problem (Akerlof, 1970)
2. **AI data governance**: Training data composition is *endogenous*, not given
   The commons degrades by design — producers' strategic opt-out decisions shape what remains
3. **Misinformation & media**: A new channel for AI to amplify misinformation
   Not by generating it, but through the systematic withdrawal of high-quality source material

Conclusion
Three contributions:
1. Economics of data markets: First evidence of adverse selection in AI training data
No price mechanism → best content exits first → lemons problem (Akerlof, 1970)
2. AI data governance: Training data composition is endogenous, not given
The commons degrades by design — producers' strategic opt-out decisions shape what remains
3. Misinformation & media: A new channel for AI to amplify misinformation
Not by generating it, but through the systematic withdrawal of high-quality source material

Let me wrap up with three contributions. First, to the economics of data markets: we provide the first large-scale evidence of adverse selection in AI training data. No price mechanism compensates quality producers, so the best content exits first—a classic lemons problem.
Second, to AI data governance: training data composition is endogenous to producers' strategic decisions. The commons degrades by design, not by accident.
Third, to misinformation research: we identify a new channel through which AI may amplify misinformation—not by generating it, but through the systematic withdrawal of high-quality source material from the training corpus.

The bottom line: as the best sources withdraw, AI training data becomes less factual, less credible, and more ideologically skewed. Without institutional solutions, this degradation becomes self-reinforcing. Thank you.

Conclusion

Three contributions:

1. **Economics of data markets**: First evidence of adverse selection in AI training data
   No price mechanism → best content exits first → lemons problem (Akerlof, 1970)
2. **AI data governance**: Training data composition is *endogenous*, not given
   The commons degrades by design — producers' strategic opt-out decisions shape what remains
3. **Misinformation & media**: A new channel for AI to amplify misinformation
   Not by generating it, but through the systematic withdrawal of high-quality source material

As the best sources withdraw, AI training data becomes systematically **less factual, less credible, and more ideologically skewed**—without institutional solutions, this degradation becomes **self-reinforcing**.

Conclusion
Three contributions:
1. Economics of data markets: First evidence of adverse selection in AI training data
No price mechanism → best content exits first → lemons problem (Akerlof, 1970)
2. AI data governance: Training data composition is endogenous, not given
The commons degrades by design — producers' strategic opt-out decisions shape what remains
3. Misinformation & media: A new channel for AI to amplify misinformation
Not by generating it, but through the systematic withdrawal of high-quality source material
As the best sources withdraw, AI training data becomes systematically less factual, less credible, and more ideologically skewed—without institutional solutions, this degradation becomes self-reinforcing.

Let me wrap up with three contributions. First, to the economics of data markets: we provide the first large-scale evidence of adverse selection in AI training data. No price mechanism compensates quality producers, so the best content exits first—a classic lemons problem.
Second, to AI data governance: training data composition is endogenous to producers' strategic decisions. The commons degrades by design, not by accident.
Third, to misinformation research: we identify a new channel through which AI may amplify misinformation—not by generating it, but through the systematic withdrawal of high-quality source material from the training corpus.
The bottom line: as the best sources withdraw, AI training data becomes less factual, less credible, and more ideologically skewed. Without institutional solutions, this degradation becomes self-reinforcing. Thank you.

Thank You!

Questions and Comments?

Kai Zhu

Bocconi University

Thank You!

Questions and Comments?

Kai Zhu

Bocconi University

Thank you. I'm happy to take questions and comments.